

9. Dynamic memory allocation

- Motivation
- *malloc/calloc*
- *free*
- How it works
- *realloc*
- Memory leaks
- Heap corruption

Motivating example

- Consider a simulation of a community of rabbits:

```
forever {
    simulationTime += deltaT;
    numRabbitsBorn = currentNumRabbits * birthRate * deltaT;
    for each new rabbit {
        construct new Rabbit record and add to population
    }
    for each rabbit in population {
        if (deathProbability * deltaT * randomInRange(0,1) > 0.5) {
            Destroy rabbit
        } else
            Have rabbit do something for its deltaT
    }
}
```

construct new Rabbit record and add to population

Destroy rabbit

How do we do these steps?

NB: The lab exercises use *Students* rather than *Rabbits* 😊



The “Rabbit Pool” solution

A rudimentary (but highly efficient) single-entity dynamic memory system.

```
typedef struct rabbit_s Rabbit;
struct rabbit_s
{
    char name[MAX_NAME_LENGTH];
    int age;
    Rabbit* next;
};

Rabbit rabbitPool[MAX_NUM_RABBITS];
Rabbit* freeRabbits = NULL;

/* Link all the rabbits in the pool onto the freeRabbits list */
void initialise(void)
{
    for (int i = 0; i < MAX_NUM_RABBITS - 1; i++) {
        rabbitPool[i].next = &rabbitPool[i + 1];
    }
    rabbitPool[MAX_NUM_RABBITS - 1].next = NULL;
    freeRabbits = &rabbitPool[0];
}
```

Continues on next slide



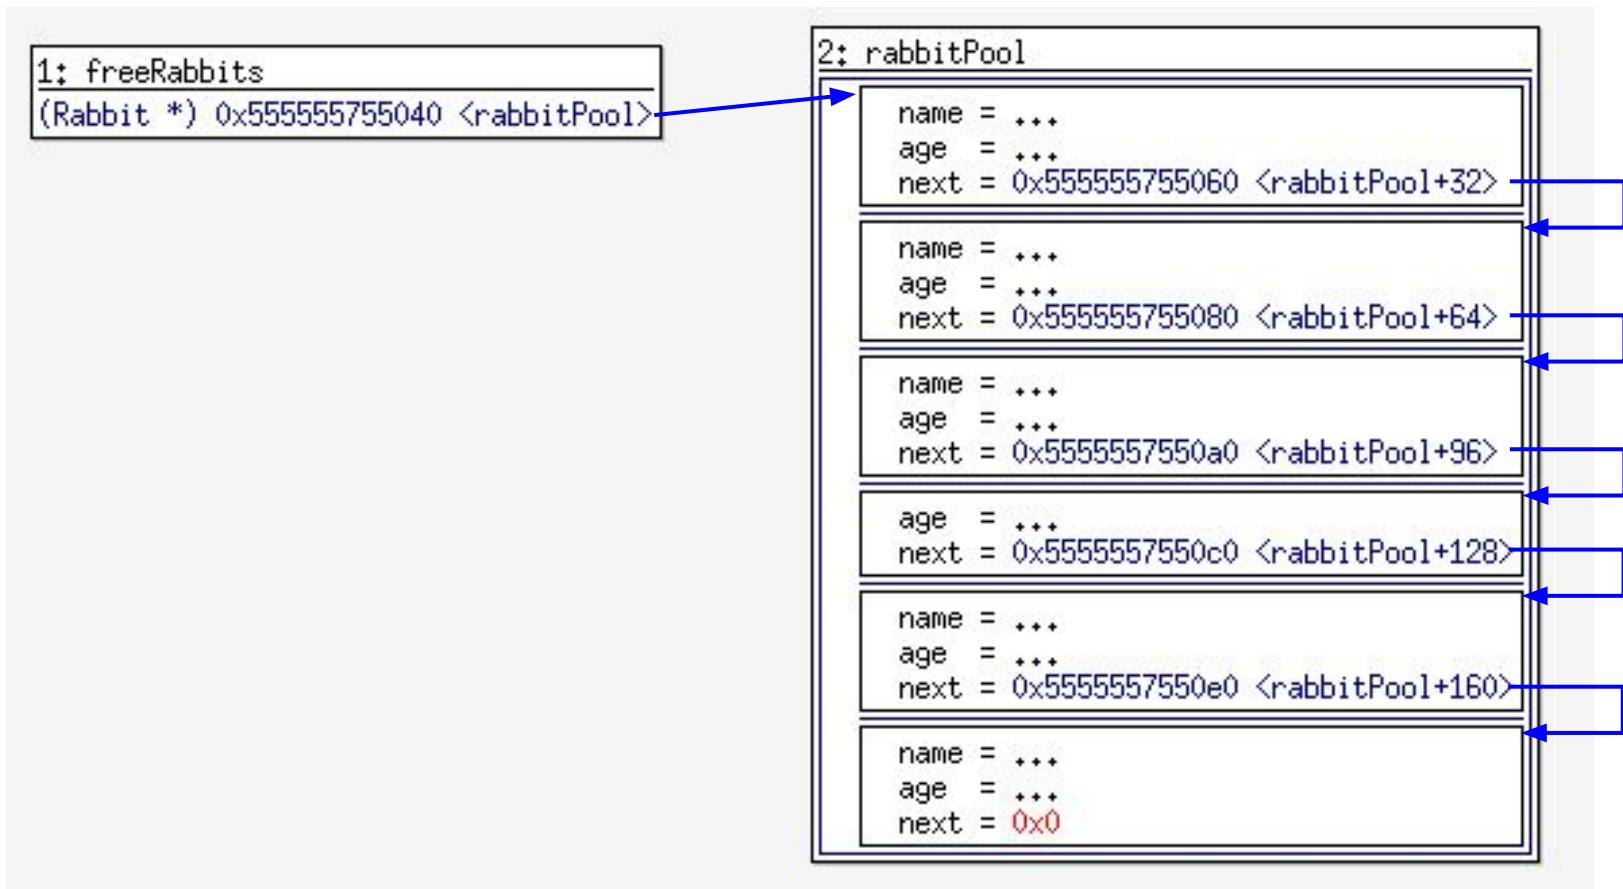
Rabbit pool (*cont'd*)

```
Rabbit* newRabbit(const char* name, int age) {
    if (freeRabbits == NULL) {
        fprintf(stderr, "Out of rabbits");
        exit(1);
    }
    Rabbit* rabbit = freeRabbits;
    strncpy(rabbit->name, name, MAX_NAME_LENGTH);
    rabbit->name[MAX_NAME_LENGTH - 1] = '\0';
    rabbit->age = age;
    freeRabbits = rabbit->next;
    rabbit->next = NULL;
    return rabbit;
}

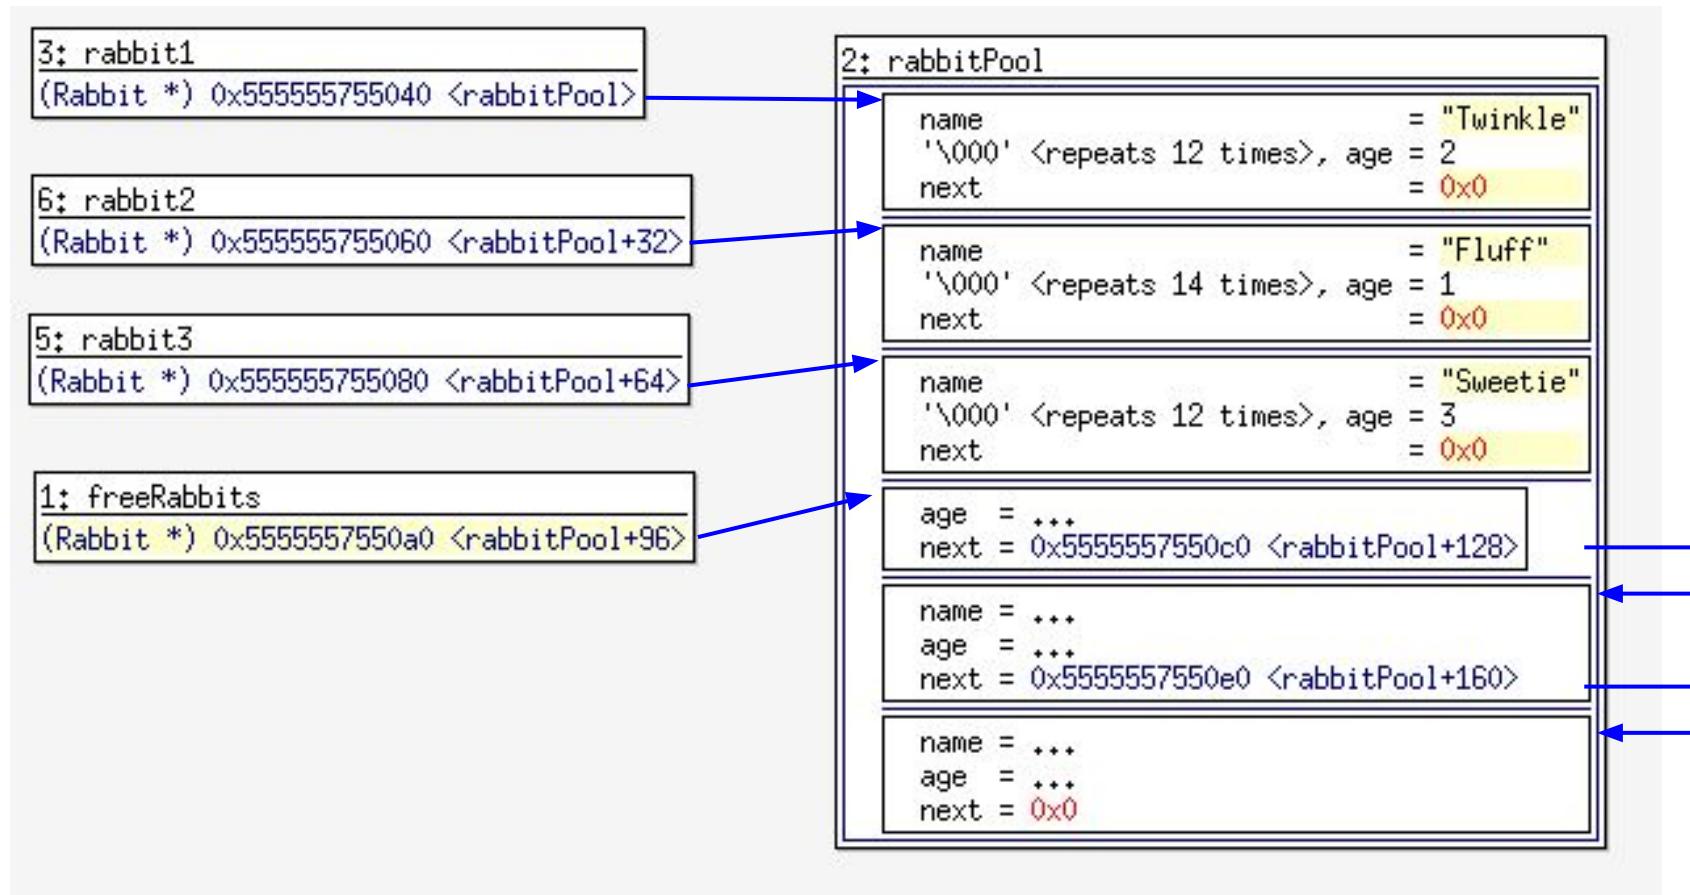
void freeRabbit(Rabbit* rabbitPtr) {
    rabbitPtr->next = freeRabbits;
    freeRabbits = rabbitPtr;
}

int main(void) {
    initialise();
    Rabbit* rabbit1 = newRabbit("Twinkle", 2);
    Rabbit* rabbit2 = newRabbit("Fluff", 1);
    Rabbit* rabbit3 = newRabbit("Sweetie", 3);
    freeRabbit(rabbit2);
}
```

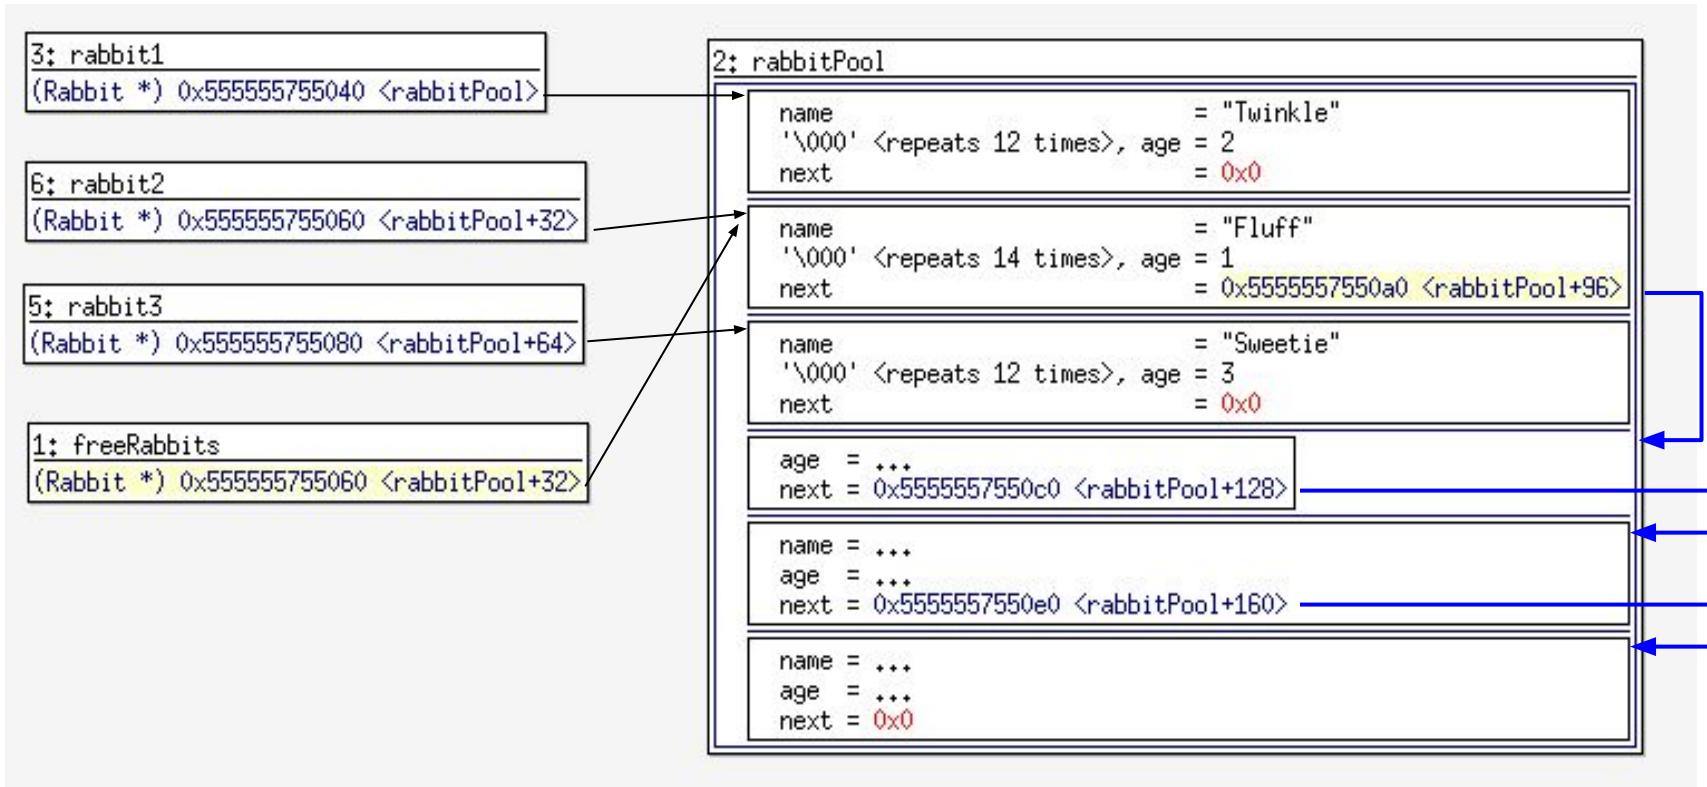
Rabbit Pool after initialise()



Rabbit Pool before freeRabbit call



Rabbit Pool after freeRabbit call



“Fluff” is back on the free list, so `*rabbit2` should not be used.



Problems

- How big is *MAX_NUM_RABBITS*?
 - If we choose it too low, program crashes
 - If we choose it too high, program wastes memory
 - May lead to unnecessary thrashing
- If we have other dynamic items to deal with (fleas, trees, grass, ...) it becomes impracticable to find good pool-size constants for all.
- SO:
 1. Allocate all new entities from a single pool of available memory
 - *malloc* and *free* library function calls
 2. Grow that pool by calls to the operating system as necessary.
 - *brk/sbrk* system calls (hidden inside *malloc*)



A modified “student” module

```
typedef struct student_s Student;

struct student_s {
    char* name;
    int age;
    Student* next;
};

// Free the student struct,
// so it can be used
void freeOneStudent(Student* sp)
{
    free(sp->name);
    free(sp);
}
```

```
Student* newStudent(const char* name,
                     int age)
{
    Student* sp = NULL;
    int nameSize = 0;
    sp = malloc(sizeof(Student));
    if (sp != NULL) {
        // if we're not out of memory
        nameSize = strlen(name);
        sp->name = malloc(nameSize + 1);
        if (sp->name == NULL) {
            // We're out of memory
            free(sp); // Free stud struct
            sp = NULL; // Must return null
        } else {
            strncpy(sp->name, name,
                    nameSize + 1);
            sp->age = age;
            sp->next = NULL;
        }
    }
    return sp;
}
```

NB!



Points to note

- *malloc* allocates a chunk of heap memory of the specified size (in bytes).
 - Returns a (`void*`) pointer (NULL if out of memory)
 - Use it wisely!
- For strings, must *malloc* (`length+1`) bytes
- *free* returns a *malloc*'d chunk of memory to the heap free list
 - Freed memory is a boojum¹
 - Looking at it will kill you.

¹*boojum* is from Lewis Carroll's poem "The hunting of the snark".

How *not* to allocate dynamic memory

- The following seems like a much easier approach.
 - And it compiles, too, if you don't use `-Wall -Werror`
- But it's CATASTROPHIC (*twice over*).
- Why?

```
Student* newStudent(const char* name, int age) {  
    Student stud;  
    stud.age = age;  
    stud.name = name; // Assume struct with char* name  
    stud.next = NULL;  
    return &stud;  
}
```

Bad!
Bad!

It's really really **IMPORTANT** that you know the answer to this (BOTH reasons).

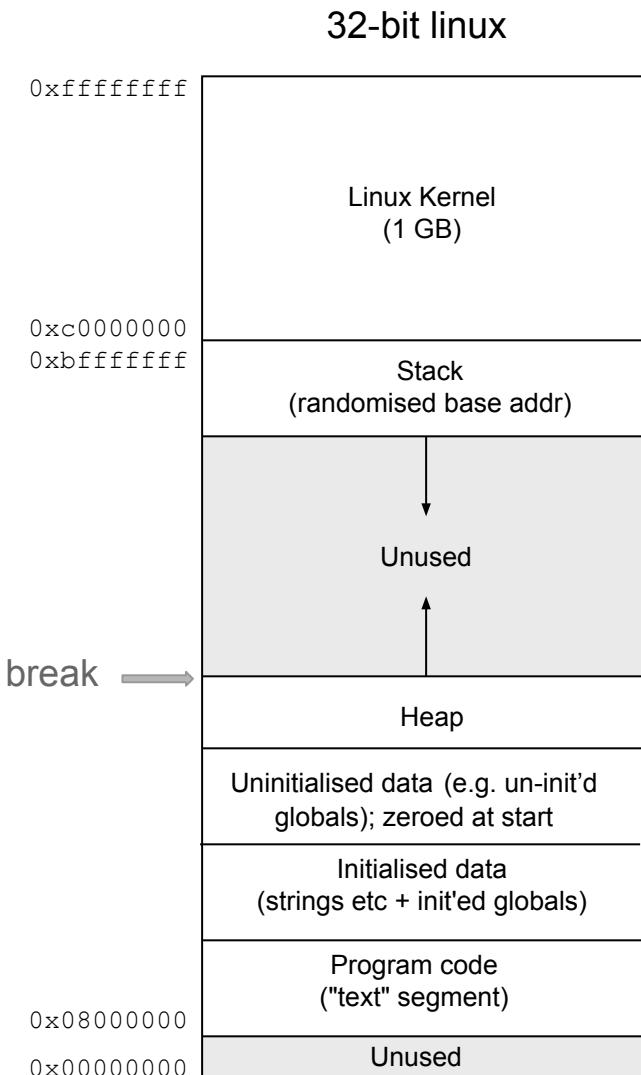


Interlude

- Request for extension on SQ2. But:
 - It's only 5% (c.f. 20% for test)
 - It's already due 2 days after end of term
 - You shouldn't spend too much time on it. The test is much more important.
- Lab question drill quiz
- Practice test available by end of week
- SQ2/Q1 & Q2 - discuss
- **Use valgrind!**

How it works

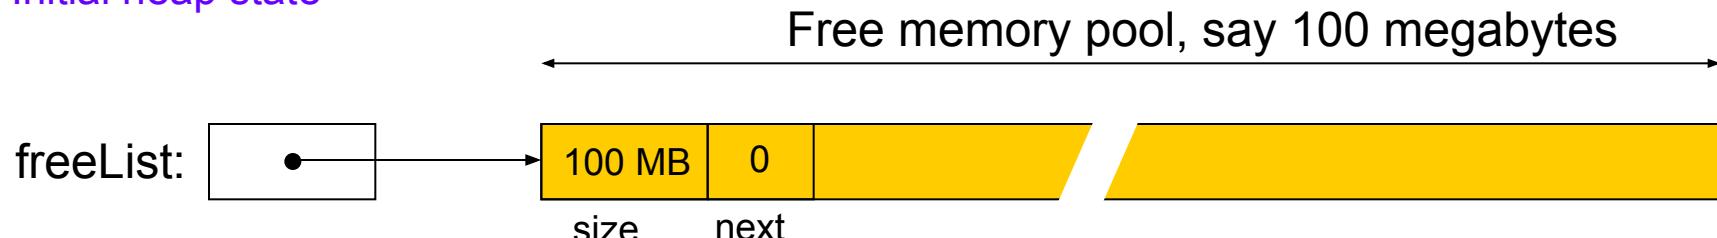
- *malloc* and *free* are the two main functions in the *memory allocator* module.
- They manage the “heap”
 - The free memory area above the initialised data segment
 - Essentially they maintain a “booking sheet” of (un)available memory
 - A global variable in the module
- The size of the heap is grown/shrunk by OS calls *brk* or *sbrk*
 - Type *man brk* to find out more



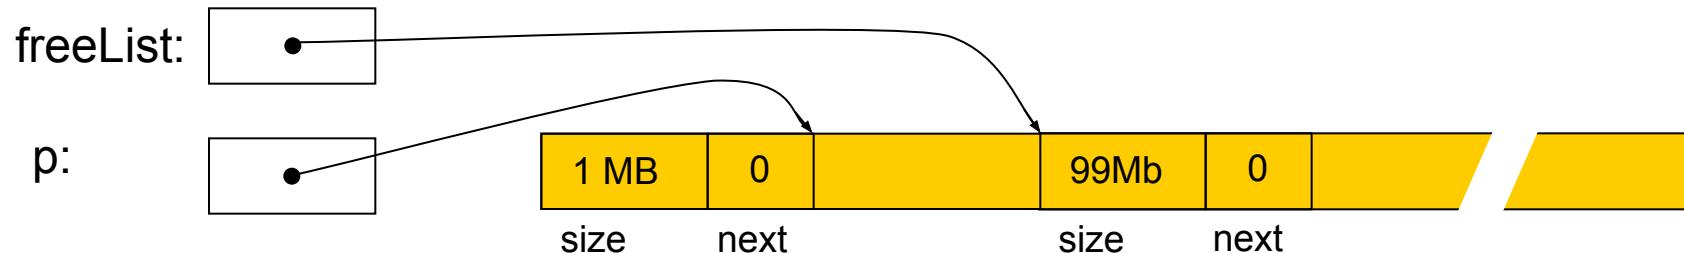
A simplistic view of malloc/free

For the full story see <https://code.woboq.org/userspace/glibc/malloc/malloc.c.html>

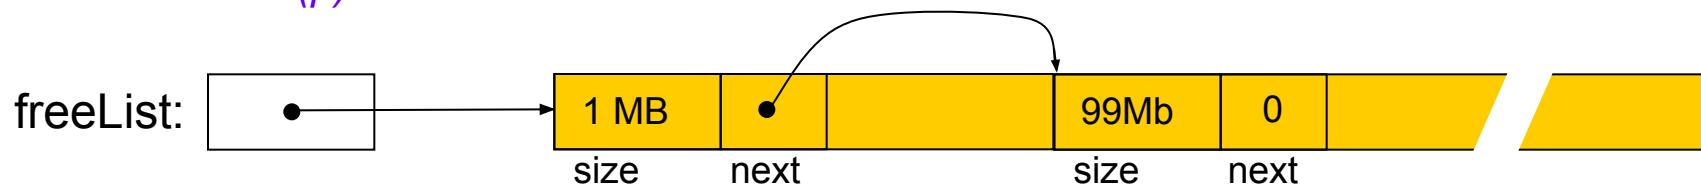
Initial heap state



State after `char *p = malloc(1000000); // 1 MB`



State after `free(p)`



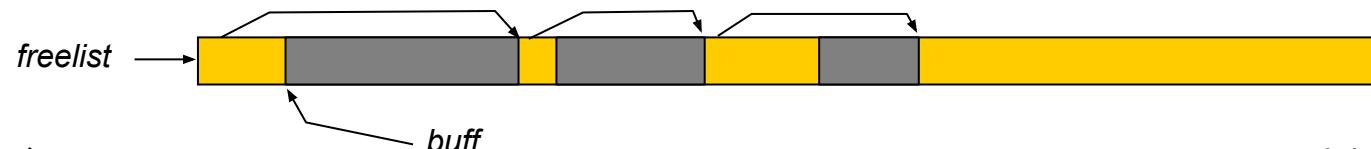
Growing malloc'd blocks

- Often need to increase size of an allocated block
- Use function `void* realloc(void* memPtr, int newSize)`
 - e.g. [Danger: bad code ahead!]

```
char* readLine(void)
{
    char* buff = NULL;
    int numBytes = 0;
    int c = 0;
    while ((c = getchar()) != EOF && c != '\n') {
        buff = realloc(buff, numBytes + 1); // Get a new bigger block
        buff[numBytes++] = c;
    }
    if (buff != NULL) {
        buff[numBytes] = '\0';
    }
    return buff; // NULL if no data read
}
```

Ptr to a malloc'd block
If *null*, realloc = malloc

NB: DON'T assume new block
is at same place as old block!





BUT ...

There are at least three things wrong with that code:

1. It has a serious bug (possibly leading to a segment fault)
 - UDOO: FIX IT!
 2. *realloc*-ing memory blocks is potentially expensive
 - may copy whole buffer if old one can't just be grown
 - doing it for every char is too inefficient for even me to tolerate
 3. The specs of that function aren't defined at the start and probably aren't what you want!
 - What are they, i.e., *exactly* what does the function return? [Consider blank lines, unterminated last lines, empty files, ...]
 - Write a better function (see *fgets* for possible spec ideas)
- Also, should be checking for *null* return from *realloc*
 - But hard to recover from it.



A better realloc strategy

- Better to *realloc* only intermittently
- Allocate an initial buffer big enough for the average case
- Double the buffer size whenever space runs out.

```
char* buff = malloc(INITIAL_BUFF_SIZE);
int bufferSize = INITIAL_BUFF_SIZE;
int numBytes = 0;
while (...) {
    if (numBytes >= bufferSize - 1) {
        bufferSize *= 2; // Double the buffer size
        buff = realloc(buff, bufferSize);
    }
    // etc
}
```

- Alternatively can increment by a fixed (largish) amount
 - But this can lead to $O(n^2)$ behaviour due to repeated buffer copying



Memory leaks

- Programs that don't free *malloc*'d memory have *memory leaks*
- Their memory footprint grows without bound
- To avoid this, every *malloc* should be matched to a corresponding *free*.
- Similarly, every call to a function like *newStudent* must have a matching call to the corresponding *freeStudent*.
- Need good clean multi-layered design
 - Match allocation and deallocation at each level



Example (from Lab 6, simplified)

```
for (i = 0; i < NUM_REPEATS; i++)  
{  
    studs = readStudents(inFile); ←  
    printStudents(&studs ); ← match  
    freeStudents(&studs ); ←  
    rewind(inputFile);  
}
```

```
StudentList readStudents(FILE* fp)  
{  
    StudentList studs = {NULL, NULL};  
    Student* sp = NULL;  
    while ((sp=readStudent(fp))!=NULL)  
    {  
        addStudent(&studs, sp); ←  
    };  
    return studs;  
}
```

```
void freeStudents(StudentList* studs)  
{  
    /* ***** TBS ***** */ ← must match  
}
```

```
Student* readStudent(FILE* fp)  
{  
    // Read from file then call ...  
    // newStudent which does ...  
    sp = malloc(sizeof(Student)); ←  
    buffSize = strlen(name) + 1;  
    sp->name = malloc(buffSize); ←  
    ...  
    strncpy(sp->name, name, buffSize);  
    sp->age = age;  
    sp->next = NULL;  
    return sp; ← match  
}  
  
void freeOneStudent(Student* sp) {  
    free(sp->name); ←  
    free(sp); ← match  
}
```

Every line of code that allocates memory should have a corresponding line that deallocates it.



Detecting memory leaks

- Need a tool for the job. Lots available. e.g.:

- *valgrind* – see <http://valgrind.org/>

- Installed on our lab machines
 - Lots of capabilities
 - Simplest:

- ```
valgrind --leak-check=yes processStudents studlist.txt
```

- Large, slow, complex internally, but fairly easy to use

- On-the-fly in-code checking with *mallinfo()*



# Heap corruption

- Over-running a *malloc*'d buffer is fatal
  - Probably
  - Eventually
- Difficult to debug
  - Solution: don't bug! *Not-bugging* is easier than *de-bugging*!
- Again there are many tools to help you find heap corruptions
  - *valgrind* checks every heap memory reference (great tool for small projects but too slow and expensive for large projects)
  - Link program with *-lmcheck*
    - Uses versions of *malloc*, *free* that do some runtime checks
    - Aborts on error
      - But checks only when *malloc*, *free* called
    - Should always use this when developing code that uses *malloc/free*



# *Parting comments on dynamic mem*

- ***Do*** understand that dynamic memory underpins all modern object-based languages
  - Objects, strings, lists etc are all allocated dynamically (i.e., on the heap)
- ***BUT***
  - ***Don't*** use dynamic memory on small microcontrollers
    - Memory is too precious, and you can't afford the risk of a memory leak
  - ***Don't*** use dynamic memory on ultra-reliable systems
    - Too hard to prove correctness
    - See [http://lars-lab.jpl.nasa.gov/JPL\\_Coding\\_Standard\\_C.pdf](http://lars-lab.jpl.nasa.gov/JPL_Coding_Standard_C.pdf)